

APR 2026

AI WORKSHOP PRODTECH OFFSITE

SURFACE → ENCODE → SELF-HEAL

90 MIN · 18 PEOPLE · 3 STATIONS · 3 ROUNDS



CONFIDENTIAL | FLOCKCOVER.COM

THREE PATTERNS

EACH BUILDS ON THE LAST. INSIGHT WITHOUT ENCODING IS A CONVERSATION. ENCODING WITHOUT SELF-HEALING IS A ONE-TIME FIX.

01

SURFACE INSIGHT

Use AI to find what's actually happening in real data. Validate the LLM's analysis. Add human judgment.

02

ENCODE FIXES

For each problem found: what's the fix, at what level? Prompt → skill → script → eval case → checklist.

03

SELF- HEALING

Design the system that catches this automatically next time. Specific enough for Claude to start building.



GOAL

THREE STATIONS

A

PRODUCT → CUSTOMER

JAY CONVOS

How does the agent that talks to fleet managers get better every week?

B

PRODUCT → ENGINEERING → CODE

TICKETS → PRs

What does developer experience look like when the agent handles planning and implementation while the engineer handles judgment?

C

PM ↔ ENG ↔ EXTERNAL

LAST MILE

The context that feeds everything — tickets, specs, docs, comms. How do we make it easier to produce good work?

THE TEAMS

EACH TEAM VISITS EVERY STATION. ANCHORS STAY AT THEIR STATION ALL THREE ROUNDS.

TEAM

1

DAVID ZAMORA

ROB GRICE

ALEX SMITH

FERGUS

TEAM

2

JORDI

JAVIER

GERAN

STE MILLINGTON

TEAM

3

SAM ADENIYI

JACOB HOLLAND

ALEKS YANEVA

ABS LAMZINI

HOW IT WORKS

THE ROTATION

ROUND 1

Teams start at their assigned station

ROUND 2

Rotate clockwise to next station

ROUND 3

Rotate clockwise again
Every team has seen every station



ANCHORS STAY AT THEIR STATION ALL THREE ROUNDS · EACH ROUND BUILDS ON WHAT THE PREVIOUS TEAM LEFT

THE TIMELINE

5 MIN

INTRO

Framework, teams, stations

20 MIN

ROUND 1 – SURFACE INSIGHT

AI analysis → human validation → categorised problem list

20 MIN

ROUND 2 – ENCODE FIXES

Each problem → fix type → specific change

20 MIN

ROUND 3 – SELF-HEALING

Design the system that catches this automatically

25 MIN

WRAP-UP

Station report-outs + synthesis

SURFACE INSIGHT

PRE-WORK DONE

Every station starts with LLM analysis of real data already complete. Nobody starts from a blank page.

YOUR JOB

Read the analysis. Test it against examples. Challenge it. What did the LLM get right? What did it miss?

THE ARTEFACT

A validated, categorised list of real problems — surfaced by AI, confirmed by humans.

THE STAKES

If this round is sloppy, everything downstream is wrong. This is the foundation.

ROUND 1 OUTPUT

✓ Category: factual errors

✓ Category: tone issues

✓ Category: tool failures

✓ Category: missed opportunities


⚠ LLM got this wrong – correction noted

⚠ Edge case: human flagged, not in AI list

VALIDATED PROBLEM LIST – HAND THIS TO ROUND 2

ENCODE FIXES

THE FIX SPECTRUM – WHERE DOES THIS CHANGE LIVE?



THE DISCIPLINE

Don't say "make it better." Write the actual change. Be specific enough that someone — or something — can implement it without asking you.

THE ARTEFACT


A table mapping each problem to a fix type and a specific change. Every row should be actionable. Vague rows don't count.

THE TEST

Could Claude implement this fix from your description alone? If not, go more specific. That's the bar.

SELF-HEALING

THE LOOP



THE QUESTION

How does this happen automatically next time? Every system generates new data. The loop is the system.

THE BAR

Write an implementation plan specific enough that Claude can start building during the wrap-up. That's "specific enough."

WHAT SEPARATES THIS

A workshop produces a list. This produces a system. The output isn't insights — it's something that runs.

A

STATION A — JAY CONVERSATIONS

DOMAIN: PRODUCT → CUSTOMER

200+ REAL JAY SESSIONS · EXTRACTED FROM DATADOG · PII-SCRUBBED

ROUND 1 SURFACE INSIGHT

Classify conversations: factual errors, hallucinations, tone issues, tool failures, user confusion, missed opportunities. What are the worst conversations? Do you trust the AI's assessment?

ROUND 2 ENCODE FIXES

For each problem category: is the fix a prompt edit, a skill update, a new tool, a deterministic script, or an eval case? Write the specific change.

ROUND 3 SELF-HEALING

Design the pipeline: every new session generates data. How does conversation → signal detection → categorisation → fix proposal → human review → deploy work?

GOAL — JAY GETS BETTER EVERY WEEK WITHOUT ANYONE MANUALLY REVIEWING EVERY CONVERSATION

B

STATION B — ENGINEERING TICKETS → PRs

DOMAIN: ENGINEERING → CODE

BATCH OF DRAFT PRs GENERATED FROM REAL LINEAR TICKETS · REVIEWED BY LLM AGAINST SOURCE TICKETS**ROUND 1 SURFACE INSIGHT**

Common failure modes across ticket definition, implementation, and code review. What categories of failure do you see? Are they in the ticket, the implementation, or the review?

ROUND 2 ENCODE FIXES

For each failure — ticket template change, CLAUDE.md update, skill, pre-flight check, automated review rule, or Notion doc? Write the specific change.

ROUND 3 SELF-HEALING

When an engineer corrects a draft PR, how does that correction persist? When a reviewer flags the same issue three times, how does it become an automated check?

C

STATION C — LAST MILE

DOMAIN: PM ↔ ENGINEERING ↔ EXTERNAL

REAL ARTEFACTS FROM THE LAST MONTH · LINEAR TICKETS · NOTION DOCS · EXTERNAL COMMS · SOME AI SLOP

ROUND 1 SURFACE INSIGHT

Rank them — which gave the recipient what they needed? Which are filler? For the bad ones: what's missing, and what's present but useless?

ROUND 2 ENCODE FIXES

For each artefact type, write a quality checklist an AI can enforce. What MUST be there? What signals slop? Model: AI critiques human writing, not the other way around.

ROUND 3 SELF-HEALING

When a ticket is created, an agent checks it against the checklist before it hits engineering. When slop gets through, how does the system learn?

GOAL — QUALITY GATES THAT IMPROVE THEMSELVES FROM EVERY ARTEFACT THAT PASSES THROUGH